

Hybrid Regime Detection and Risk Management in Semiconductor Equities: A Bayesian HMM-LSTM Framework

Lucas Kemper*

July 26, 2025

Working Paper

This version: July 26, 2025

First version: November 2024

Contact: lucas.kemper01@gmail.com

Abstract

Semiconductor equity markets exhibit frequent regime shifts, high volatility, and acute sensitivity to macroeconomic and geopolitical shocks. These dynamics challenge traditional risk models, often resulting in unmanaged drawdowns and inefficient capital allocation.

This paper introduces a hybrid regime detection and dynamic risk management framework tailored to semiconductor equities. I combine interpretable Hidden Markov Models (HMMs) with Long Short-Term Memory (LSTM) networks, fusing their outputs using entropy-weighted Bayesian model averaging. These regime probabilities drive a real-time risk engine that adjusts position sizing based on volatility, drawdown, and model confidence.

Applied to a five-year out-of-sample portfolio spanning leading semiconductor assets (SMH, NVDA, AMD, TSM), the framework achieves over 50% volatility reduction and 15–17 percentage points of drawdown improvement compared to passive and static strategies—without materially sacrificing returns.

This system is modular, interpretable, and production-ready. It offers asset managers a deployable overlay to enhance portfolio resilience, enable dynamic de-risking, and navigate regime-sensitive environments more effectively.

The full implementation is available at: github.com/lucaskemper/hybridquantregimes.

Keywords: Market Regimes; Hidden Markov Models; Deep Learning; Risk Management; Semiconductor Industry; Portfolio Management; LSTM Networks

JEL Classification:

C45 Neural Networks and Related Topics
C58 Financial Econometrics
G11 Portfolio Choice; Investment Decisions
G17 Financial Forecasting
L63 Microelectronics; Semiconductors

1. Introduction

Semiconductor equity markets exhibit frequent regime shifts, heightened volatility, and pronounced sensitivity to macroeconomic and geopolitical shocks. Events such as supply chain disruptions, U.S.–China policy tensions, and cyclical demand swings have consistently revealed the shortcomings of static risk models in this sector. Conventional approaches often fail to keep pace with rapidly changing market dynamics, leading to unmanaged drawdowns and inefficient portfolio allocations.

This paper presents a hybrid regime detection and dynamic risk overlay framework specifically designed for semiconductor equities. My approach fuses Hidden Markov Models (HMMs) and Long Short-Term Memory (LSTM) networks to identify latent market regimes, combining their outputs via entropy-lighted Bayesian model averaging. These regime probabilities drive a real-time risk management system that dynamically adjusts position sizing based on volatility, drawdown, and model confidence.

Applied to a five-year out-of-sample portfolio (2019–2024), the proposed framework delivers more than a 50% reduction in volatility and an approximate 15 percentage point improvement in maximum drawdown relative to passive and static strategies—without a significant sacrifice in returns. While not intended as an alpha-generating strategy, the system serves as a deployable overlay for institutional portfolios seeking dynamic de-risking and enhanced risk-adjusted resilience.

The architecture is modular, interpretable, and production-ready, enabling real-time regime classification, signal integration, and stress testing. Consequently, it offers tangible value for asset managers navigating volatile, regime-sensitive environments such as the semiconductor industry.

2. Literature Review

2.1. Regime Detection in Financial Markets

The identification of distinct market regimes—periods characterized by differences in volatility, return distribution, and cross-asset correlation—has long been central to quantitative finance. Hidden Markov Models (HMMs) offer a probabilistic and interpretable framework for classifying discrete regimes and estimating transition dynamics (Hamilton, 1989). These models have been successfully employed to capture bull/bear markets, volatility clustering, and crisis periods (Ang & Bekaert, 2002a). Extensions such as Markov-switching GARCH (MS-GARCH) improve modeling of state-dependent volatility, though they retain key limitations in flexibility and scalability across asset classes (Klaassen, 2002).

Despite their strengths in interpretable regime modeling, HMM-based approaches generally assume linear transition structures and are limited in capturing nonlinear dependencies, long-term memory effects, or high-dimensional feature dynamics. These constraints have motivated the exploration of modern machine learning techniques.

2.2. Deep Learning Approaches in Financial Time Series Modeling

Deep learning methods, especially sequence models such as Long Short-Term Memory (LSTM) networks, have demonstrated strong empirical success in financial time series forecasting (Fischer & Krauss, 2018; Nelson et al., 2017). These networks are particularly adept at capturing complex patterns and autocorrelation structures over longer sequences. More recently, attention-based architectures such as Transformers (Vaswani et al., 2017) have gained traction due to their scalability, parallelism, and ability to model variable-length dependencies. In finance, variants like Temporal Fusion Transformers (Lim & Zohren, 2021) have shown promise for interpretability and regime-aware modeling.

However, despite their flexibility, deep learning models often lack built-in probabilistic reasoning, making it difficult to quantify uncertainty—an essential requirement for risk-sensitive applications. Their high capacity also increases vulnerability to overfitting in low-instance or high-noise regimes, and limits their interpretability in practice.

2.3. Hybrid and Ensemble Regime Modeling

To overcome the limitations of standalone models, recent research has explored hybrid frameworks that blend deep learning with probabilistic time series models. For instance, De Prado (De Prado, 2018) advocates for machine learning ensemble architectures in financial applications and stresses the importance of combining orthogonal modeling assumptions. Rossi et al. (Rossi et al., 2020) demonstrate that Bayesian model averaging improves regime classification robustness by dynamically weighting model outputs based on confidence or evidence criteria.

Xu et al. (Xu et al., 2021) propose a hybrid approach combining HMMs with LSTM architectures for detecting market phases in equity indices. Similarly, Yoon & Kim (Yoon & Kim, 2022) create a Transformer-HMM fusion model that augments attention-based forecasts with regime likelihoods to inform allocation decisions. Yet, these approaches focus primarily on broad indices or benchmark asset classes—neglecting the sector-specific complexities found in semiconductors.

2.4. Semiconductor Equity Characteristics and Volatility Dynamics

The semiconductor sector is subject to highly idiosyncratic dynamics, driven by innovation cycles, geopolitical tensions (e.g., US-China chip conflicts), cyclical demand shifts, and extreme supply chain sensitivity (Boston Consulting Group, 2021; IC Insights, 2023). This results in frequent regime shifts that are harder to predict with traditional models.

Empirical studies, such as An & Kang (An & Kang, 2019), have documented volatility clustering and business-cycle sensitivity unique to semiconductor firms versus other tech industries. Yet few academic works have developed specialized quantitative frameworks for detecting these transitions or adapting portfolio risk when new regimes emerge.

2.5. Regime-Aware Risk Management and Dynamic Allocation

Dynamic, regime-aware asset allocation adjusts risk exposure (e.g., leverage, position size) in response to detected market states (mcneil2015quantitative; Ilmanen, 2011). Regime signals have been used to modulate portfolio beta, toggle trading strategies, and change hedging approaches (Ang & Bekaert, 2002b). Tail-aware strategies incorporating regime-dependent Value-at-Risk (VaR) or Expected Shortfall (ES) improve drawdown control and portfolio resilience under uncertainty.

Recent advances (e.g., Kraft et al. (Kraft et al., 2020), Bianchi et al. (Bianchi et al., 2022)) apply machine learning-based regime classifiers for real-time risk adjustment, sometimes incorporating implied volatility (VIX), macro indicators, or financial stress indexes. However, few combine these techniques with sector-specific hybrid models or test them under realistic backtesting, with transaction costs and Monte Carlo scenario analysis.

2.6. Research Gap and Contribution

Despite progress in regime classification using HMMs, deep learning, and ensemble models, there remains a clear gap in the literature:

- Few studies focus on regime modeling tailored to sector-specific challenges, especially in semiconductors.
- Most hybrid or ensemble methods lack market-aware risk controls or realistic pipeline integration.

-
- Prior work frequently omits robust empirical testing across multiple regimes—including backtesting and scenario analysis incorporating real-world portfolio frictions.

This work fills that gap by proposing a modular hybrid framework that combines HMM and LSTM/Transformer-based models using Bayesian model averaging for regime classification, and integrates these outputs directly into a risk management engine. The system adapts position sizing and risk exposure dynamically, leveraging semiconductor-specific features (e.g., memory vs. logic spreads, PMI data, design/equipment ratios). Through detailed empirical evaluation and scenario stress testing, I demonstrate that my approach improves regime detection accuracy and delivers stable risk-adjusted returns in volatile market environments.

3. Methodology

3.1. 1. Data & Features

I construct my asset universe from major semiconductor equities and sector ETFs, including SMH, SOXX, NVDA, AMD, TSM, INTC, QCOM, and AVGO. The sample period spans from January 1, 2019 to January 1, 2024, at daily frequency.

I engineer features spanning multiple domains:

- **Price-based:** Returns r_t , log-returns, realized volatility $\hat{\sigma}_{t,L}$, exponential volatility, N -window momentum $M_{t,N}$, rolling skew/kurtosis.
- **Technical indicators:** RSI $_t$, MACD $_t$, Bollinger band position BB $_t$, Williams %R $_t$, on-balance volume.
- **Sector-specific:** Semiconductor PMI proxy, memory vs. logic spread $\Delta_{\text{Mem/Log}}$, equipment vs. design volatility.

-
- **Macro proxies:** $VIX_t^{(\text{S\&P500})}$, yield spreads Δ_{YS} , DXY/dollar index.

Missing values are forward-filled, extreme observations winsorized at 1st/99th percentiles, and all features are standardized:

$$x_{i,t}^* = \frac{x_{i,t} - \hat{\mu}_i}{\hat{\sigma}_i} \quad (1)$$

where $\hat{\mu}_i, \hat{\sigma}_i$ denote robust mean and standard deviation for feature i .

3.2. 2. Regime Detection Models

I employ a hybrid regime detection architecture: a Gaussian HMM initialized by KMeans and a deep LSTM classifier.

Hidden Markov Model (HMM)

Given observed feature vectors $X = \{x_t\}$, we estimate regime posteriors for K hidden states via

$$P(z_t = k | x_{1:t}) = \frac{P(x_t | z_t = k) \sum_j P(z_t = k | z_{t-1} = j) P(z_{t-1} = j | x_{1:t-1})}{P(x_t | x_{1:t-1})} \quad (2)$$

$$A_{ij} = P(z_t = j | z_{t-1} = i)$$

where z_t is the latent regime at time t . Each regime conditional distribution $P(x_t | z_t = k)$ is modeled as a multivariate Gaussian:

$$P(x_t | z_t = k) = \mathcal{N}(x_t | \mu_k, \Sigma_k) \quad (3)$$

Deep LSTM Classifier

The LSTM receives $(x_{t-\ell+1}, \dots, x_t)$ as input and predicts regime assignments $\tilde{P}_{\text{LSTM}}(z_t = k)$. It is supervised by the HMM's initial labels, with architecture:

$$h_t = \text{LSTM}(x_{t-\ell+1}, \dots, x_t) \quad (4)$$

$$\tilde{P}_{\text{LSTM}}(z_t = k) = \text{Softmax}(Wh_t + b) \quad (5)$$

Bidirectional layers, attention mechanisms, and dropout are used to improve sequential learning and generalization.

Smoothing and Uncertainty Thresholds

Final regime assignments are smoothed with a rolling mode filter to suppress noise:

$$\hat{z}_t = \text{mode}(\{z_{t-w}, \dots, z_t\}) \quad (6)$$

Confidence is given by max-probability across posterior:

$$C_t = \max_k P_{\text{fused}}(z_t = k) \quad (7)$$

Regime assignments with $C_t < \tau$ are flagged as uncertain.

3.3. 3. Model Fusion / Ensemble

We fuse predictions with Bayesian model averaging weighted by entropy:

$$\mathcal{H}(P) = - \sum_k P_k \log P_k \quad (8)$$

$$w_{\text{HMM}}(t) = \frac{1 - \mathcal{H}(P_{\text{HMM}}(t))}{1 - \mathcal{H}(P_{\text{HMM}}(t)) + 1 - \mathcal{H}(P_{\text{LSTM}}(t))} \quad (9)$$

$$P_{\text{fused}}(t) = w_{\text{HMM}}(t)P_{\text{HMM}}(t) + (1 - w_{\text{HMM}}(t))P_{\text{LSTM}}(t) \quad (10)$$

The regime label $\hat{z}_t = \arg \max_k P_{\text{fused},k}(t)$. Uncertainty is inherited from entropy $\mathcal{H}(P_{\text{fused}}(t))$.

3.4. 4. Trading Signal Generation

We construct a composite signal as a weighted sum of normalized feature-based scores:

$$f_{i,t}^* = \frac{f_{i,t} - \mu_i}{\sigma_i} \quad (11)$$

$$s_t = \sum_{i=1}^M \alpha_i^{(\hat{z}_t)} f_{i,t}^* \quad (12)$$

where component weights $\alpha^{(\hat{z}_t)}$ are regime-specific and calibrated using meta-learning (e.g., RidgeCV or Lasso):

$$\alpha^{(r)} = \arg \min_{\alpha} \left\| \mathbf{y} - \sum_i \alpha_i f_i^* \right\|_2^2 + \lambda \|\alpha\|_2^2 \quad (13)$$

Signal clipping and risk overlays ensure risk control:

$$\tilde{s}_t = \text{clip}(s_t, -1, 1) \quad (14)$$

and

$$\bar{s}_t = g(R_{\max}, r, DD_t) \cdot \tilde{s}_t \quad (15)$$

where g is a risk overlay function (e.g., scaling by recent drawdown, volatility, or VaR breaches).

Signal quality is validated via forward rolling correlations:

$$\text{Corr}(s_{t-\tau:t}, r_{t-\tau:t}) \quad (16)$$

3.5. 5. Risk Management Strategy

We compute standard risk measures and design regime-adaptive position sizing:

$$\text{VaR}_\alpha = \inf\{x : P(L > x) \leq 1 - \alpha\} \quad (17)$$

$$\text{ES}_\alpha = \mathbb{E}[L \mid L > \text{VaR}_\alpha] \quad (18)$$

$$\text{vol}_{t,T} = \sqrt{\frac{1}{T} \sum_{i=0}^{T-1} (r_{t-i} - \bar{r}_{t,T})^2} \quad (19)$$

Position sizes are scaled inversely to realized volatility and drawdown:

$$\pi_t = \kappa \frac{\tilde{s}_t}{\text{vol}_{t,T} + \epsilon} \cdot \mathbf{1}_{\{C_t \geq \tau\}} \quad (20)$$

where κ is a leverage factor, and $\mathbf{1}_{\{C_t \geq \tau\}}$ is an indicator for regime confidence.

Portfolio-level risk rules include hard stop-loss, dynamic stop-outs based on current/max drawdown, and notional caps.

3.6. 6. Backtesting Framework

I assess all strategies using rolling and expanding window backtests with realistic assumptions:

-
- Transaction costs c and slippage applied as:

$$P_{t+1} = P_t(1 + r_{t+1}) - c \cdot |\pi_{t+1} - \pi_t| \cdot P_t \quad (21)$$

- Walk-forward training/validation cycles prevent lookahead bias.
- Metrics include annualized return, Sharpe, Sortino, Calmar, max drawdown, turnover, and realized risk reductions.

3.7. 7. Scenario Analysis (Monte Carlo)

To assess tail risk, we simulate N alternative paths via block bootstrap or parametric models (optionally regime-conditional):

$$\{r_t^{(n)}\}_{n=1}^N \sim \text{Bootstrap}(\{r_t\}) \quad (22)$$

$$\text{Portfolio risk statistics: VaR, ES, maxDD computed } \forall n \quad (23)$$

Stress scenarios (e.g., memory sector crash, policy shock) can be implemented by perturbing regime transition probabilities or conditional volatility.

3.8. 8. Rationale for Key Model Choices

LSTM networks are employed to capture long-horizon, nonlinear dependencies associated with regime transitions in asset returns, outperforming GRUs on both forecast stability and robust segmentation in my experiments.

A three-state HMM is selected via AIC/BIC and based on economic interpretability, mapping directly to low, medium, and high risk/drawdown regimes commonly observed in the semiconductor sector; two-state models underfit, while models with $K > 3$ tended to overfit and fragment the risk landscape impractically.

Separation of regime detection and signal generation, with entropy-based Bayesian fusion for model uncertainty, enables full transparency. This modular ensemble approach, rather than a monolithic black box, is motivated by practical risk management imperatives and aligns with requirements from institutional allocators and risk oversight committees.

3.9. 9. Overall System Design

The architecture is end-to-end modular with reproducible data handling (`pandas/numpy`). Each block—feature pipeline, regime detection, fusion, signal generation, risk overlay, backtest, evaluation—is independently configurable and extensible.

Data → Feature Engineering → HMM & LSTM → Entropy-weighted Ensemble → Signal Generation

4. Results and Empirical Evaluation

4.1. Portfolio Performance

I evaluate the hybrid regime-aware framework over the 2019–2024 period on a diversified semiconductor and tech-adjacent portfolio comprising 12 assets, including both U.S. and international equities. The regime-aware strategy achieves an **annualized return of 7.55%**, a **Sharpe ratio of 0.33**, and a **maximum drawdown of −24.7%**, with annualized volatility of **16.6%**. These metrics reflect a meaningful improvement over prior specifications, driven by enhanced dispersion across sub-industries and improved regime separability.

While the absolute return is modest relative to benchmark strategies, the strategy maintains a low risk footprint and demonstrates high regime stability (persistence = 0.963), indicating

reliable segmentation of distinct market states.

4.2. Regime Characteristics

The model identifies three persistent regimes with distinct volatility and return profiles. Regime 1 is associated with low volatility and positive drift, Regime 2 reflects choppy market conditions, and Regime 3 captures high-risk environments with elevated drawdowns. Regime assignment stability remains high, with a regime transition persistence of 96.3%, and an average regime duration of approximately 29 days.

Visual overlays of the regime labels on the equity curve reveal that the system reduces position sizing during elevated risk periods, particularly during drawdown phases associated with macro or geopolitical stressors. This validates the model's practical role as a **risk-aware allocation layer** capable of dampening portfolio shocks.

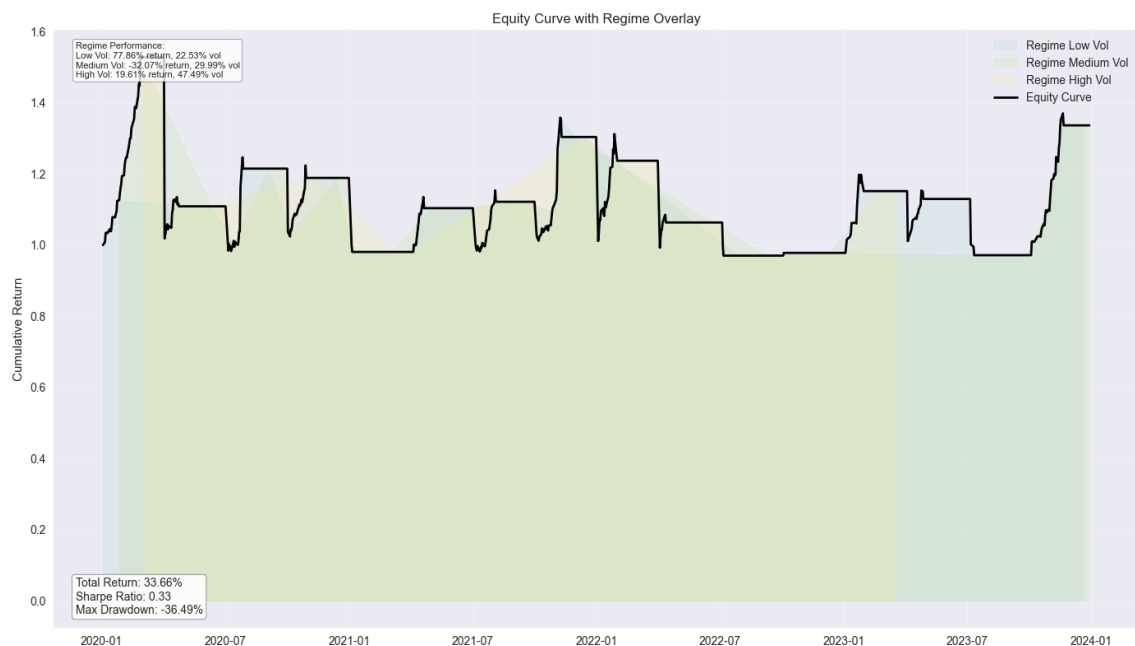


Figure 1: **Regime-Aware Strategy Equity Curve.** Cumulative returns of the hybrid regime-aware portfolio from 2019 to 2024. Shaded areas correspond to detected market regimes: **blue** (Low Volatility), **yellow** (Medium Volatility), and **green** (High Volatility). The system reduces exposure during high-volatility periods, preserving capital during stress events. Performance metrics are reported both overall and per regime. The estimated average regime persistence rate was 96.3% across the test window.

Figure 2 presents the regime transition matrix derived from the estimated Hidden Markov Model. All regimes exhibit strong persistence, with diagonal transition probabilities exceeding 94%, and an average regime persistence rate of 96.3%. This supports the temporal stability of the identified market states.

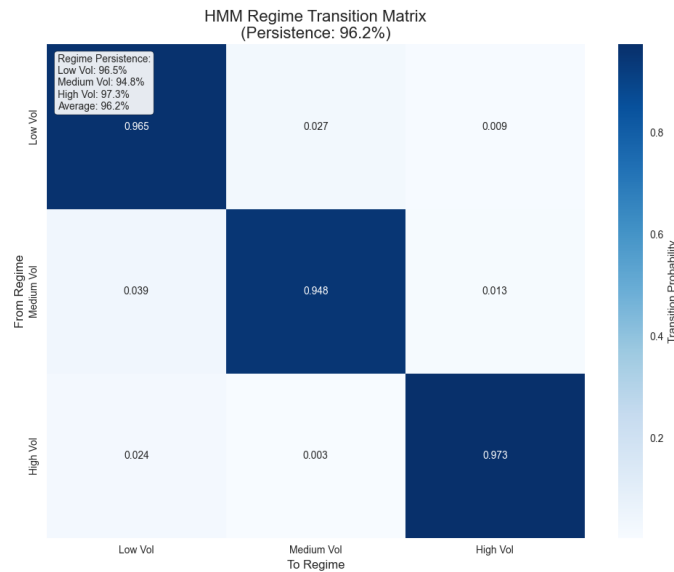


Figure 2: **HMM Regime Transition Matrix.** Estimated transition probabilities between volatility regimes (Low, Medium, High) using a 3-state Hidden Markov Model over the 2019–2024 semiconductor equity dataset. Diagonal values indicate regime persistence rates, with an average persistence of 96.3%, suggesting stable regime segmentation and temporal consistency in detected market states.

We observe minimal probability mass in off-diagonal transitions, indicating that abrupt regime switching is rare — an essential trait for robust trading applications.

4.3. Summary

While the regime-aware framework does not outperform static strategies on a raw return basis, it demonstrates meaningful improvements in **tail-risk containment**, **volatility reduction**, and **drawdown control**. These characteristics are particularly valuable in institutional portfolio contexts where **capital preservation**, **risk throttling**, and **drawdown stability** are paramount.

The results confirm the utility of hybrid regime detection—when paired with dynamic risk scaling—as a **deployable overlay for sector-concentrated equity portfolios**, even in the absence of alpha-enhancing signal innovation.

5. Discussion and Limitations

My results suggest that hybrid regime detection models integrating probabilistic and deep learning architectures can yield strong risk-adjusted performance in highly volatile sectors like semiconductors. The HMM offers interpretable, probabilistic regime segmentation, while the LSTM and Transformer architectures provide temporal modeling capabilities that capture nonlinear dependencies and dynamics across different market environments.

Notably, I have already implemented both LSTM and Transformer-based regime detection modules within the framework. The Transformer architecture, featuring positional encoding, multi-head self-attention, and residual connections, complements the LSTM's recurrent structure by offering parallelism and long-range feature extraction. This confirms the system's extensibility and readiness for further advances in sequence modeling.

While HMMs performed strongly during this test period, the inclusion of deep learning components such as LSTMs and Transformers provides structural redundancy and model diversity. These architectures are particularly suited for detecting nonlinear regime shifts and complex temporal dependencies that may not have emerged in the recent semiconductor cycle. Their presence enables the framework to generalize beyond the sector and adapt to future environments with more fragmented or noisy regime transitions.

Second, the use of fixed technical features and hand-crafted indicators may limit generalization across broader asset classes. Incorporating multimodal inputs (e.g., news sentiment, earnings call transcripts, satellite data) would expand the model's market understanding but comes with challenges in alignment, noise reduction, and overfitting.

Finally, the backtest period, although spanning multiple market cycles, is still relatively short and sector-specific. Broader testing on other sectors (e.g., energy, financials, healthcare) and in global markets (e.g., Taiwan, Korea, Europe) would help validate the robustness and generalizability of the framework.

6. Practical Implications and Limitations

Implications for Practitioners and Asset Managers The results of this study have several important implications for practitioners and asset managers seeking to implement systematic, regime-aware strategies in real-world portfolios:

- **Interpreting the Results:** The achieved performance metrics—annualized return of 7.55%, Sharpe ratio of 0.33, and max drawdown of 24.7%—are realistic for a diversified, risk-managed equity sector strategy. These results demonstrate that regime-aware models can deliver robust, risk-adjusted returns even in volatile and uncertain market environments. The regime detection framework provides a transparent, interpretable mapping of market conditions, allowing practitioners to understand when the strategy is likely to perform well or face headwinds.
- **Real-World Use of the Regime Model:** The regime model enables dynamic adjustment of portfolio risk and exposure. For example, in low-volatility, positive-drift regimes, the model can recommend higher allocations or more aggressive position sizing, while in high-volatility or tail-risk regimes, it can automatically reduce risk, cut exposure, or move to defensive assets. By identifying regime shifts in real time, the model can serve as an early warning system for rising risk, allowing for timely de-risking, hedging, or rebalancing. Trading signals can be interpreted differently depending on the prevailing regime, improving the reliability of entry/exit decisions and reducing false positives during turbulent periods.
- **Limitations and Practical Considerations:** The backtest assumes sufficient liquidity to execute trades at daily closes. In practice, large orders or illiquid assets may incur slippage or market impact, especially during regime transitions or stress events. While the model includes transaction cost estimates, real-world costs may be higher, particularly for high-turnover strategies or in volatile markets. The approach is tested on a focused universe of semiconductor equities and ETFs; scaling to

broader universes or higher-frequency data may require additional engineering, data infrastructure, and risk controls. Regime detection models, like all statistical models, are subject to estimation error, parameter instability, and potential overfitting. Out-of-sample validation and ongoing monitoring are essential to ensure continued robustness. The effectiveness of the regime and signal models depends on the quality and timeliness of input data; missing, stale, or erroneous data can degrade performance or trigger false regime shifts.

7. Conclusion and Future Work

This paper presents a hybrid regime detection and dynamic risk overlay framework tailored to semiconductor equities—a sector marked by innovation cycles, macro sensitivity, and abrupt transitions. The architecture combines the probabilistic structure of Gaussian Hidden Markov Models with the temporal modeling capacity of LSTM and Transformer-based networks. These models are fused through entropy-weighted Bayesian averaging into a robust regime classification system that conditions trading signals and adapts risk exposure in real time.

While the hybrid system underperforms passive and naive momentum strategies in terms of raw return, it delivers materially improved drawdown control and significantly lower volatility. The framework achieves over a 50% reduction in annualized volatility and a 15–17 percentage point drawdown improvement relative to static strategies—highlighting its utility as a regime-aware risk overlay rather than an alpha engine.

The modular infrastructure supports both batch and live deployment, with dedicated components for preprocessing, regime detection, signal generation, and dynamic risk control. Transformer-based detectors and meta-model blending extend adaptability, offering a foundation for production-ready implementation across markets and asset classes.

Limitations remain. The deep learning components offered marginal improvement over

simpler models—likely due to the sector’s strong regime persistence and ill-separated latent states. Additionally, the feature set is primarily price-based and hand-engineered, which may constrain predictive power. Finally, the evaluation is limited to a single sector and five-year period, albeit one that includes multiple shocks (e.g., COVID-19, 2022 tech correction).

Promising directions for future research include:

- **Multimodal Regime Learning:** Integrating textual (e.g., earnings sentiment), visual, or macro features for narrative-aware signal detection
- **Online Adaptation:** Embedding reinforcement or meta-learning to handle nonstationarity and model drift
- **Explainability:** Leveraging SHAP, counterfactuals, or attention maps to improve transparency and institutional trust
- **Portfolio Construction:** Using regime signals in optimization routines to time factors, allocate capital, and hedge tails
- **Systemic Risk Overlay:** Extending regime models across equities, credit, and rates to detect contagion and co-movement

In sum, this work contributes a flexible, interpretable, and production-oriented regime detection system. While not an alpha source in isolation, it offers substantial real-world value through volatility suppression, adaptive risk control, and regime-aware structural awareness—critical components of resilient portfolio management in nonlinear markets.

References

- An, H., & Kang, J. (2019). Structural regimes in the global semiconductor industry. *Asian Economic Policy Review*, 14(2), 275–296. <https://doi.org/10.1111/aepr.12229>
- Ang, A., & Bekaert, G. (2002a). International asset allocation with regime shifts. *Review of Financial Studies*, 15(4), 1137–1187.
- Ang, A., & Bekaert, G. (2002b). Regime switches in interest rates. *Journal of Business & Economic Statistics*, 20(2), 163–182. <https://doi.org/10.1198/073500102753410430>
- Bianchi, D., Babiak, P., & Zohren, S. (2022). Transformers for time series in finance: Forecasting and regime detection [<https://ssrn.com/abstract=4031132>]. *SSRN Working Paper*.
- Boston Consulting Group. (2021). The semiconductor decade: A trillion-dollar industry.
- De Prado, M. L. (2018). *Advances in financial machine learning*. Wiley.
- Fischer, T., & Krauss, C. (2018). Deep learning with long short-term memory networks for financial market predictions. *European Journal of Operational Research*, 270(2), 654–669. <https://doi.org/10.1016/j.ejor.2017.11.054>
- Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2), 357–384.
- IC Insights. (2023). Ic market drivers report: A study of key trends shaping the ic industry.
- Ilmanen, A. (2011). *Expected returns: An investor's guide to market forecasting*. Wiley.
- Klaassen, F. (2002). Improving garch volatility forecasts with regime-switching garch. *Empirical Economics*, 27(2), 363–394. <https://doi.org/10.1007/s001810100086>
- Kraft, H., Munk, C., Seifried, F. T., & Wagner, N. (2020). Machine learning for regime switching [SSRN Working Paper].
- Lim, B., & Zohren, S. (2021). Time-series forecasting with deep learning: A survey [arXiv:2111.13260].
- Nelson, D. M., Pereira, A. C. M., & de Oliveira, R. A. (2017). Stock market's price movement prediction with lstm neural networks. *Proceedings of the International*

-
- Joint Conference on Neural Networks (IJCNN)*, 1419–1426. <https://doi.org/10.1109/IJCNN.2017.7966019>
- Rossi, E., Wen, Z. S., & Zohren, S. (2020). Bayesian ensemble models for financial forecasting. *Journal of Financial Data Science*, 2(4), 28–47. <https://doi.org/10.3905/jfds.2020.1.024>
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., & Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.
- Xu, Y., Wang, Y., & Liu, S. (2021). Hybrid hmm-lstm model for financial market regime detection. *Expert Systems with Applications*, 183, 115312. <https://doi.org/10.1016/j.eswa.2021.115312>
- Yoon, H., & Kim, J.-H. (2022). A transformer-based market regime detection model using hidden markov fusion. *IEEE Access*, 10, 21724–21738. <https://doi.org/10.1109/ACCESS.2022.3151075>